MPCS 51087
~~Problem Set 4~~
GPU Ray Tracing with CUDA

~~Summer 2019~~

~~Due Date: Thursday, August 15, 2019 by 11:59 pm~~

1 Introduction to Ray Tracing

Ray tracing is a powerful method for rendering three-dimensional objects with complex light interactions. Many other physical phenomena are also analogous to ray tracing – for example, simulating neutrons passing through matter. Figure 1 shows the basic idea for a reflective object. An observer (the eyeball) is viewing an object through a window (the rectangle). The object is illuminated by a light source, which is emitting rays of light (the arrows) in many directions. The observer will see rays that reflect off of the object.


Figure 1: Illustration of Ray-Tracing. Image credit: [1]

Our task is to render the image seen by the observer through the window. To do so, the simulation can do one or both of the following:

- Simulate light rays starting at the light source and ending at the observer.

- Simulate light rays starting from the observer and going backwards to the light source.

We will implement the latter scheme in this assignment. We will simulate a reflective sphere in black-and-white, illuminated by a single light source. A serial version can be implemented in about 100 lines of code.
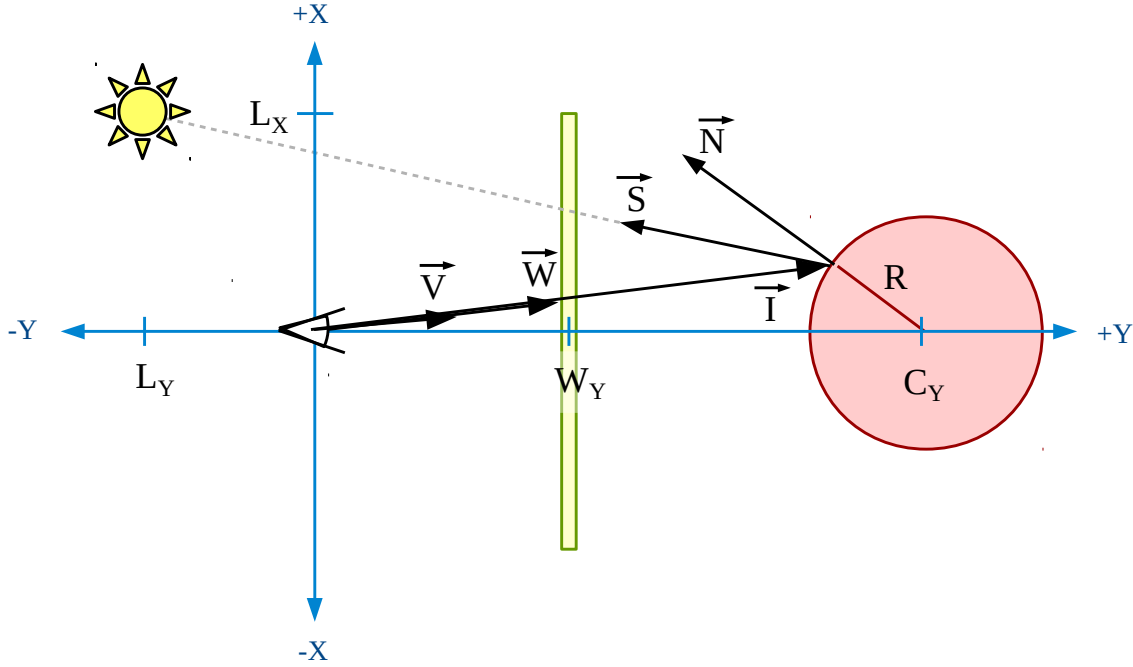
Figure 2: 2D Diagram for Ray-Tracing

2   Vector Notation

Solving the problem involves the use of 3D vectors. A vector $\vec{V}$ has scalar components, $\vec{V} = (V_X, V_Y, V_Z)$. Adding or subtracting two vectors yields a vector, $\vec{V} - \vec{U} = (V_X - U_X, V_Y - U_Y, V_Z - U_Z)$. Multiplying a scalar and a vector yields a scalar, $t\vec{V} = (tV_X, tV_Y, tV_Z)$. Dividing a vector by a scalar yields a scalar, $\vec{V}/t = (V_X/t, V_Y/t, V_Z/t)$. The dot product (a multiplication between two vectors) yields a scalar, $\vec{V} \cdot \vec{U} = V_X U_X + V_Y U_Y + V_Z U_Z$. The norm (or magnitude or length) of the vector is a scalar defined by $|\vec{V}| = \sqrt{V_X^2 + V_Y^2 + V_Z^2}$.

3   Theory

This section describes the theory used in our implementation. The implementation itself is brief and is described in the algorithms below, so you may skip this section as needed.

Figure 2 shows a 2D cross-section of the problem with the information we need. The observer is at the origin and faces the positive y-direction. The sphere is located at $\vec{C} = (C_X, C_Y, C_Z)$ and has radius R. The window is parallel to the (x,z)-plane at $y = W_Y$ and has bounds $-W_{max} < W_X < W_{max}$ and $-W_{max} < W_Z < W_{max}$. In the simulation, we will represent the window as an $n \times n$ grid, G. The light source is located at $\vec{L} = (L_X, L_Y, L_Z)$.

Our task is to simulate many rays originating from the observer (so-called "view rays") with randomly-selected directions. The steps are as follows:

- **Select the direction of view ray ($\vec{V}$).** Let $\vec{V}$ be a unit vector representing the direction of the view ray. In spherical coordinates, we will randomly select its component angles (θ, ϕ).

Then we will get $\vec{V}$ in Cartesian coordinates:

$$V_X = \sin\theta\cos\phi$$
$$V_Y = \sin\theta\sin\phi$$
$$V_Z = \cos\theta$$

One important point to note is that we cannot sample both ϕ and θ from uniform random distributions between 0 and 2π without biasing the distribution. To sample the unit sphere correctly, we will sample only ϕ from a uniform distribution between 0 and 2π. We will then sample the *cosine* of θ from a uniform distribution between -1.0 and 1.0. E.g.,

---
Algorithm 1 Direction Sampling Algorithm

1: Sample ϕ from uniform distribution $(0, 2\pi)$
2: Sample $\cos(\theta)$ from uniform distribution $(-1, 1)$
3: $\sin(\theta) = \sqrt{1 - \cos^2(\theta)}$
4: $V_X = \sin(\theta)\cos(\phi)$
5: $V_Y = \sin(\theta)\sin(\phi)$
6: $V_Z = \cos(\theta)$

Also note that we cannot sample V_X, V_Y, and V_Z directly from a uniform distribution $(-1, 1)$ in all three Cartesian dimensions and then normalize, as this would generate a biased distribution.

- **Find the intersection of the view ray with the window ($\vec{W}$).** Knowing that the window is at W_Y, the window's point-of-intersection with the view ray is given by the vector $\vec{W}$:

$$\vec{W} = \frac{W_Y}{V_Y}\vec{V}$$

If the view ray is outside the window ($|W_X| > W_{max}$ or $|W_Z| > W_{max}$), we reject it and chose a new $\vec{V}$.

- **Find the intersection of view ray with sphere ($\vec{I}$).** Let $\vec{I}$ be the sphere's point-of-intersection with the view ray. To find $\vec{I}$, we solve the following system of equations:

$$\vec{I} = t\vec{V}$$
$$\left|\vec{I} - \vec{C}\right|^2 = R^2$$

These are the equations of the view ray and the sphere, respectively. Solving for t yields:

$$t = \left(\vec{V} \cdot \vec{C}\right) - \sqrt{\left(\vec{V} \cdot \vec{C}\right)^2 + R^2 - \vec{C} \cdot \vec{C}}$$

which can be back-substituted to get $\vec{I}$. If t does not have a real solution $((\vec{V} \cdot \vec{C})^2 + R^2 - \vec{C} \cdot \vec{C} < 0)$, then view ray does not intersect the sphere and we choose a new $\vec{V}$.

- **Find the observed brightness of the sphere (b).** Next, we want to find the brightness of the sphere that is observed at $\vec{I}$. To do so , we:

  - *Find the unit normal vector ($\vec{N}$).* The unit normal vector $\vec{N}$ is perpendicular to the sphere's surface at $\vec{I}$.

$$\vec{N} = \frac{\vec{I} - \vec{C}}{\left|\vec{I} - \vec{C}\right|}$$

   – *Find the direction to the light source ($\vec{S}$).* The direction to light source (sometimes called the "shadow ray") is represented by the unit vector $\vec{S}$.

$$\vec{S} = \frac{\vec{L} - \vec{I}}{\left| \vec{L} - \vec{I} \right|}$$

   – *Find the brightness (b).* The brightness can be found from $\vec{S}$ and $\vec{N}$ using "Lambertian shading".

$$b = \left\{ \begin{array}{ll} 0 & \vec{S} \cdot \vec{N} < 0 \\ \vec{S} \cdot \vec{N} & \vec{S} \cdot \vec{N} \geq 0 \end{array} \right.$$

- **Add the brightness to the window's grid.** We find (i, j) such that $\vec{G}(i, j)$ is the position of $\vec{W}$ on the window's grid G and let:

$$G(i, j) = G(i, j) + b$$

4 Implementation

Algorithm 2 describes a ray-tracing implementation. As described above, the observer is at the origin and is facing the positive-y direction. The sphere is located at $\vec{C} = (C_X, C_Y, C_Z)$ and has radius R. The light source is located at $\vec{L} = (L_X, L_Y, L_Z)$. The window is parallel to the (x,z)-plane at $y = W_Y$, has bounds $-W_{max} < W_X < W_{max}$ and $-W_{max} < W_Z < W_{max}$, and is represented by an $n \times n$ grid, G.

Algorithm 2 Ray Tracing Algorithm

1: allocate $G[1 \ldots n][1 \ldots n]$ $\triangleright$ The window is represented on the grid G
2: $G[i][j] = 0$ for all (i, j)
3: **for** $n = 1 \ldots N_{rays}$ **do**
4:      **repeat**
5:          Sample random $\vec{V}$ from unit sphere $\triangleright$ Algorithm 1
6:          $\vec{W} = \frac{W_Y}{V_Y} \vec{V}$ $\triangleright$ The intersection of the view ray and the window
7:      **until** $|W_X| < W_{max}$ **and** $|W_Z| < W_{max}$ **and** $\left(\vec{V} \cdot \vec{C} \right)^2 + R^2 - \vec{C} \cdot \vec{C} > 0$
8:      $t = \left(\vec{V} \cdot \vec{C} \right) - \sqrt{\left(\vec{V} \cdot \vec{C} \right)^2 + R^2 - \vec{C} \cdot \vec{C}}$
9:      $\vec{I} = t \vec{V}$ $\triangleright$ The intersection of the view ray and the sphere
10:     $\vec{N} = \frac{\vec{I} - \vec{C}}{|\vec{I} - \vec{C}|}$ $\triangleright$ The unit normal vector at $\vec{I}$
11:     $\vec{S} = \frac{\vec{L} - \vec{I}}{|\vec{L} - \vec{I}|}$ $\triangleright$ The direction of the light source at $\vec{I}$
12:     $b = \text{MAX} \left(0, \vec{S} \cdot \vec{N} \right)$ $\triangleright$ The brightness observed at $\vec{I}$
13:     find (i, j) such that $G(i, j)$ is the gridpoint of $\vec{W}$ on G
14:     $G(i, j) = G(i, j) + b$ $\triangleright$ Add brightness to grid (must be thread safe!)
15: Output grid G to file

    Note that in Algorithm 2, all rays are independent from one another, except for the fact that they are all tallying their results to a single common grid (Line 14). As the rays are randomized, it may be the case that multiple rays try to update the value at the same grid location at the same

time, resulting in undefined behavior. Therefore, some sort of synchronization will be necessary to ensure correctness.

Figure 3 shows a sample image where $\vec{L} = (4, 4, -1)$, $W_Y = 10$, $W_{max} = 10$, $C = (0, 12, 0)$, $R = 6$.
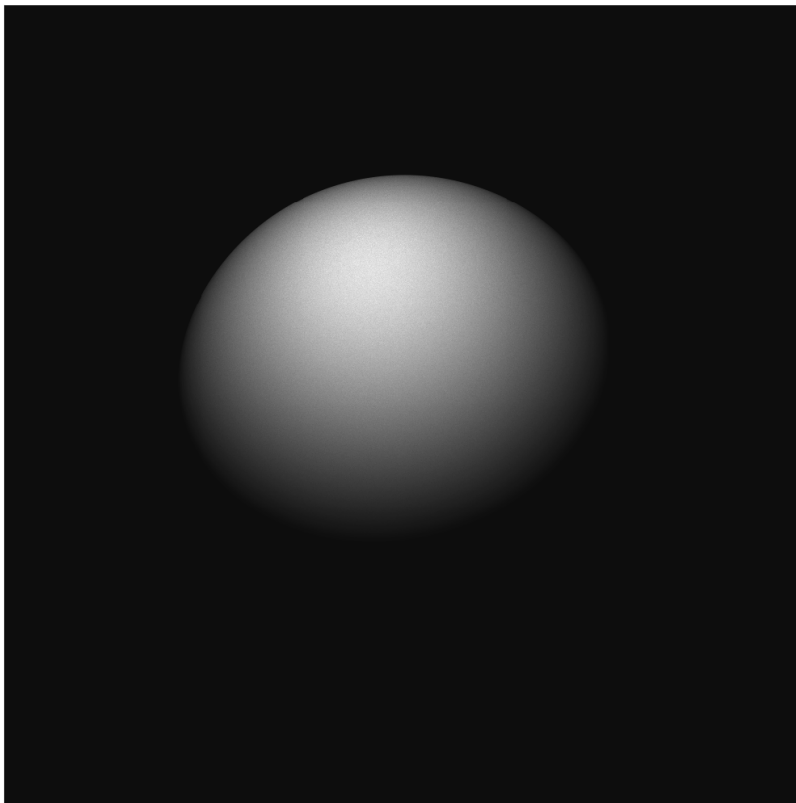

Figure 3: Ray-Traced Render of a Sphere Illuminated from Top Left

4.1   Random Number Generation on GPU

There are a wide variety of methods for generating pseudo random number sequences on computers. The simplest methods are known as Linear Congruential Generators (LCGs), and are very fast and only require a singe variable (typically a 64-bit unsigned integer) to track the state of the system. As such, they are very efficient in terms of speed and memory usage, though are subject to several limitations. One limitation is that some generators have a small period, where a period is the quantity of numbers that it can generate before the same number gets generated again and the cycle begins to repeat itself. For instance, the standard C function `rand()` has a period of roughly 2^{32}, or roughly 2 billion. Another limitation is that LCGs typically have low *k-dimensionality*, meaning that if sampling a multi-dimensional object (say, points in 3D space), the points end up being highly correlated and tend to form hyperplanes rather than be truly uniform in space. A final weakness of LCGs is that the numbers they generate tend to have less randomness in their low order bits, meaning that some parts of the numbers are in a sense more random than others. This can be particularly troublesome if say a random integer is being used in a modulus operation where only the low order bits are being used.

There are more robust methods such as Mersenne Twister and PCG algorithms. These algorithms

can potentially provide much higher periods, higher k-dimensionality, and better statistical quality for all generated bits. In the case of the Mersenne Twister algorithm, this is accomplished by having a much larger state vector, composed of hundreds of numbers, that work with the algorithm to define where the generator is in the stream.

On GPUs, the standard CUDA method of generator random numbers is by use of the `curand` functions. This is an implementation of the Mersenne Twister algorithm that has a large state vector, which can cause some issues if each of millions of threads needs to initialize a unique generator stream to just generate a handful of random numbers. Typically, this will mean that it will be necessary to reduce the total number of curand random number generator streams that are created to be far below the number of threads being launched, or otherwise make re-use of each random number stream. For instance, if you want to launch 1 billion threads, it is best to either reduce the number of threads by packing multiple tasks into each thread, or set up a sharing system for the RNG streams.

However, `curand` is not the only option for generating numbers on the GPU! We can also just use our own algorithm that we implement ourselves. For this assignment, you will have the following two options for generating random numbers:

1. Use the `curand` functions, but make sure that each thread samples and traces multiple rays, so as to amortize the costs of initializing the generator and reducing the overall memory cost of the generator. Be sure to follow appropriate procedures for ensuring different threads generate different random values.

2. -OR- Launch each ray on its own thread, and use an LCG pseudo random number generator (PRNG) below:

```
// A 63−bit LCG
// Returns a double precision value from a uniform distribution
// between 0.0 and 1.0 using a caller−owned state variable.
double LCG_random_double(uint64_t * seed)
{
  const uint64_t m = 9223372036854775808ULL; // 2^63
  const uint64_t a = 2806196910506780709ULL;
  const uint64_t c = 1ULL;
  *seed = (a * (*seed) + c) % m;
  return (double) (*seed) / (double) m;
}
```

If using an LCG PRNG, you will need a way to ensure that each thread has a uniquely initialized state (or seed) variable to pass to the generator so that the generator's output is unique to that thread. One way to do this is to set the seed to some function of the thread's index (e.g., `seed = thread_id ×4238811`), though a more robust method is to point each thread to a different location within a PRNG's period by "forwarding". If we assume each ray will require at most 100 direction samples to generate a valid direction that passes the test on line 7 of Algorithm 2, and each direction sample requires two uses of the PRNG, then if we assign one thread per ray we can simply forward each thread N times, where $N = 200×$`thread_id`. Forwarding can in theory by done naively by repeatedly calling the LCG algorithm N times, but thankfully a much more efficient $log(N)$ method is available by using modular exponentiation:

```
1  // "Fast Forwards" an LCG PRNG stream
2  // seed: starting seed
3  // n: number of iterations (samples) to forward
4  // Returns: forwarded seed value
5  uint64_t fast_forward_LCG(uint64_t seed, uint64_t n)
6  {
7    const uint64_t m = 9223372036854775808ULL; // 2^63
8    uint64_t a = 2806196910506780709ULL;
9    uint64_t c = 1ULL;
10   n = n % m;
11   uint64_t a_new = 1;
12   uint64_t c_new = 0;
13   while(n > 0)
14   {
15     if(n & 1)
16     {
17       a_new *= a;
18       c_new = c_new * a + c;
19     }
20     c *= (a + 1);
21     a *= a;
22     n >>= 1;
23   }
24   return (a_new * seed + c_new) % m;
25 }
```

5 Questions

1. (20 pts) Make and test a serial CPU version of our ray tracing code by implementing Algorithm 2. Take user input for the number of rays and gridpoints; the other parameters can be hardcoded. (Hint: Define structs and functions for the vectors and vector operations. With these simple abstractions, the code can be written in about 100 lines.)

2. (30 pts) Use CUDA to parallelize your code to run on GPUs. Note that, while the rays are independent, updating the window could result in write conflicts. Therefore, some synchronization (e.g., atomic operations) will be necessary. Your code should run almost entirely on GPU, with the completed grid being copied back to CPU only at the end for outputting to file by the host.

3. (10 pts) Experiment with different block and thread configurations of your code and determine an (approximately) optimal configuration for launching your kernel on a Midway GPU of your choice. Report your optimal settings in your PDF writeup.

4. (20 pts) Compare the runtime of the serial CPU version and CUDA GPU versions as functions of problem size (the number of rays). For the CUDA version, use your optimal block/thread configuration developed from the previous problem.

5. (10 pts) Show a sample image of at least 1000×1000 grid resolution produced by your CUDA ray-tracing implementation generated using at least 1 billion rays.

6 Compiling, Code Cleanliness, and Documentation (10 pts)

We plan on compiling and running your code as part of grading it. You must include a makefile capable of compiling your code into a runnable executable. Also, while we are not grading to any rigid coding standard, your code should be human readable and well commented so we can understand what is going on.

7 What to Submit

You submission should be in the form of a zip or tar file containing the following items:

- Your source code for your serial and CUDA versions should be in different directories, each with their own makefiles.

- Working makefiles that compile each version of your code

- Your plotting script(s)

- A single writeup (in PDF form) containing all plots and discussions asked for in the assignment

8 References

1. Rademacher, Paul. *Ray Tracing: Graphics for the Masses.* `https://www.cs.unc.edu/~rademach/xroads-RT/RTarticle.html`

2. *Ray tracing (graphics).* `http://en.wikipedia.org/wiki/Ray_tracing_(graphics)`